


Tachyon: A Reliable Memory Centric Storage for Big Data Analytics

Haoyuan (HY) Li, Ali Ghodsi, Matei Zaharia,
Scott Shenker, Ion Stoica



TACHYON

Outline



- Overview
- Research
- Open Source
- Future

Outline

- **Overview**
- Research
- Open Source
- Future

Memory is King

- RAM throughput increasing exponentially
- Disk throughput increasing slowly



Memory-locality key to interactive response time

Realized by many...

- Frameworks already leverage memory



April 7, 2012

Many kinds of memory-centric data management

I'm frequently asked to generalize in some way about in-memory or memory-centric data management. I can start:


- The desire for [human real-time interactive response](#) naturally leads to



Problem solved?

An Example: Spark


- Fast in-memory data processing framework
 - Keep **one** in-memory copy inside JVM
 - Track **lineage** of operations used to derive data
 - Upon failure, use lineage to recompute data



Issue 1

*Different jobs share data:
Slow writes to disk*


storage engine &
execution engine
same process
(slow writes)



Issue 1


*Different frameworks share data:
Slow writes to disk*

storage engine &
execution engine
same process
(slow writes)




Issue 2

execution engine &
storage engine
same process



Issue 2


execution engine &
storage engine
same process



Issue 2


Process crash: lose all cache

execution engine &
storage engine
same process



Issue 3

duplicate memory per job & GC



Tachyon

Reliable data sharing at ***memory-speed***
within and across cluster frameworks/jobs

Solution Overview

Basic idea

- Push **lineage** down to storage layer
- Use memory aggressively

Facts

- One data copy in memory
- Rely on recomputation for fault-tolerance


Stack

Computation Frameworks
(Spark, MapReduce, Impala, Tez, ...)


Tachyon

Existing Storage Systems
(HDFS, S3, GlusterFS, ...)

Architecture




Lineage



Issue 1 revisited


*Different frameworks share
at memory-speed*

execution engine &
storage engine
same process
(fast writes)




Issue 2 revisited

execution engine &
storage engine
same process



Issue 2 revisited


execution engine &
storage engine
same process



Issue 2 revisited

process crash: keep memory-cache


execution engine &
storage engine
same process



Issue 3 revisited

*Off-heap memory storage
one memory copy & no GC*

execution engine &
storage engine
same process
(no duplication & GC)



Outline

- Overview
- Research
- Open Source
- Future

Question 1: How long to get missing data back?




That server contains
the data computed
last month!



Lineage enables **Asynchronous Checkpointing**

Edge Algorithm

- Checkpoint leaves
- Checkpoint hot files
- Bounded Recovery Cost



Question 2: How to allocate recomputation resource?




Would recomputation
slow down my high
priority jobs?
Priority Inversion?




Recomputation Resource Allocation


- Priority Based Scheduler




- Fair Sharing Based Scheduler




Comparison with in Memory HDFS




Further Improve Spark's Performance



Grep





Master Faster Recovery



Recomputation Resource Consumption

Bin	Tasks	% of Jobs	
		Facebook	Bing
1	1 - 10	85%	43%
2	11 - 50	4%	8%
3	51 - 150	8%	24%
4	151 - 500	2%	23%
5	> 500	1%	2%

Trace Summary



Outline

- Overview
- Research
- Open Source
- Future




TACHYON Open Source Status

- Apache License, Version 0.4.1 (Feb 2014)

PUBLIC  [amplab / tachyon](#)  [Unstar](#) 813

- 15+ Companies
- Spark and MapReduce applications can run without any code change

Release Growth



Contributors Outside of Berkeley

More than 75%

Tachyon is the
Default Off-Heap Storage
Solution for  Spark

Tachyon
is in Fedora 20

Thanks to Redhat!


Commercially
Supported
By





TACHYON in *Spark* *H₂O*

Today, data gets parsed and exchanged between Spark and H2O via Tachyon. Users can interactively query big data both via SQL and ML from within the same context.



Spark/MapReduce/Shark without Tachyon

- Spark
 - `val file = sc.textFile("hdfs://ip:port/path")`
- Hadoop MapReduce
 - `hadoop jar hadoop-examples-1.0.4.jar wordcount hdfs://localhost:19998/input hdfs://localhost:19998/output`
- Shark
 - `CREATE TABLE orders_cached AS SELECT * FROM orders;`

Spark/MapReduce/Shark with Tachyon

- Spark
 - val file = sc.textFile("tachyon://ip:port/path")
- Hadoop MapReduce
 - hadoop jar hadoop-examples-1.0.4.jar wordcount tachyon://localhost:19998/input tachyon://localhost:19998/output
- Shark
 - CREATE TABLE orders_tachyon AS SELECT * FROM orders;

Spark OFF_HEAP with Tachyon

```
// Input data from Tachyon's Memory  
val file = sc.textFile("tachyon://ip:port/path")
```

```
// Store RDD OFF_HEAP in Tachyon's Memory  
file.persist(OFF_HEAP)
```

Thanks to our Code Contributors!

Aaron Davidson

Achal Soni

Ali Ghodsi

Andrew Ash

Anurag Khandelwal

Aslan Bekirov

Bill Zhao

Calvin Jia

Colin Patrick McCabe

Shivaram Venkataraman

Chang Cheng

Du Li

Fei Wang

Gerald Zhang

Grace Huang

Hao Cheng

Haoyuan Li

Henry Saputra

Hobin Yoon

Huamin Chen

Jey Kottalam

Joseph Tang

Lukasz Jastrzebski

Manu Goyal

Mark Hamstra

Nick Lanham

Orcun Simsek

Pengfei Xuan

Qifan Pu

Qianhao Dong

Raymond Liu

Reynold Xin

Robert Metzger

Rong Gu

Sean Zhong

Srinivas Parayya

Tao Wang

Timothy St. Clair

Vamsi Chitters

Xi Liu

Xiaomin Zhang




Outline

- Overview
- Research
- Open Source
- Future

Goal?

Better Assist Other Components



Welcome Collaboration!

Short Term Roadmap

- Ceph Integration (Redhat)
- Hierarchical Local Storage (Intel)
- Further improver Shark Performance (Yahoo)
- Better support for Multi-tenancy (AMPLab)
- ***Many more*** from AMPLab and Industry Collaborators.

Your Requirements?

Tachyon Summary

- High-throughput, fault-tolerant memory centric storage, with lineage as a first class citizen
- Further improve performance for frameworks such as Spark, Hadoop, and Shark etc.
- Healthy community with 15+ companies contributing

Thanks!

Questions?

- *More Information:*
 - <http://tachyon-project.org>
 - <https://github.com/amplab/tachyon>
- *Email:*
 - haoyuan@cs.berkeley.edu