

Graph *Flow*

Using Spark and ~~Shark~~ SparkSQL to
Power a Real-time Recommendation
and Customer Intelligence Platform

About

- @MLnick
- Co-founder @graphflow - big data & machine learning applied to recommendations, consumer behaviour & insights
- Apache Spark committer
- Author of “Machine Learning with Spark”
 - *Packt RAW: <http://www.packtpub.com/machine-learning-with-spark/book>*

Graph *Flow*


Agenda

- A bit about Graphflow
- Architecture overview:
 - *Graphflow Platform*
 - *Why Spark?*
 - *Evolution of our architecture*
- Future

Graph *Flow*

1

Graphflow's platform captures interactions between users and items:



Items can be almost anything, from **products to movies to music to apps**.

Interactions can include **view, purchase, click, like** and anything in between.

2


We use state-of-the art **machine learning** and **large-scale analytics techniques** to analyse this behaviour and generate recommendations for users and items.

Our flexible integration layer can deliver **real-time, personalised** recommendations across any channel, including **online, mobile, email** and **SMS**.

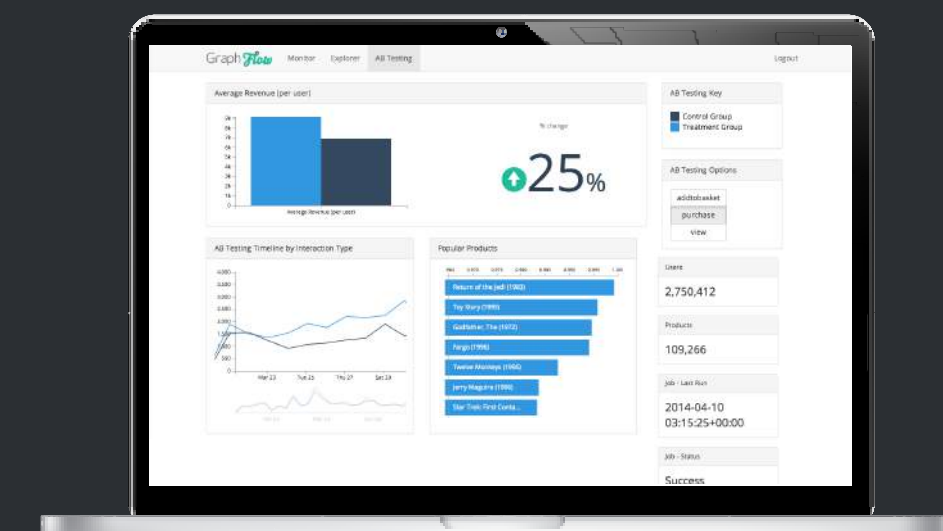
Our analytics portal provides **deep insight** into customer behaviour and product performance

3

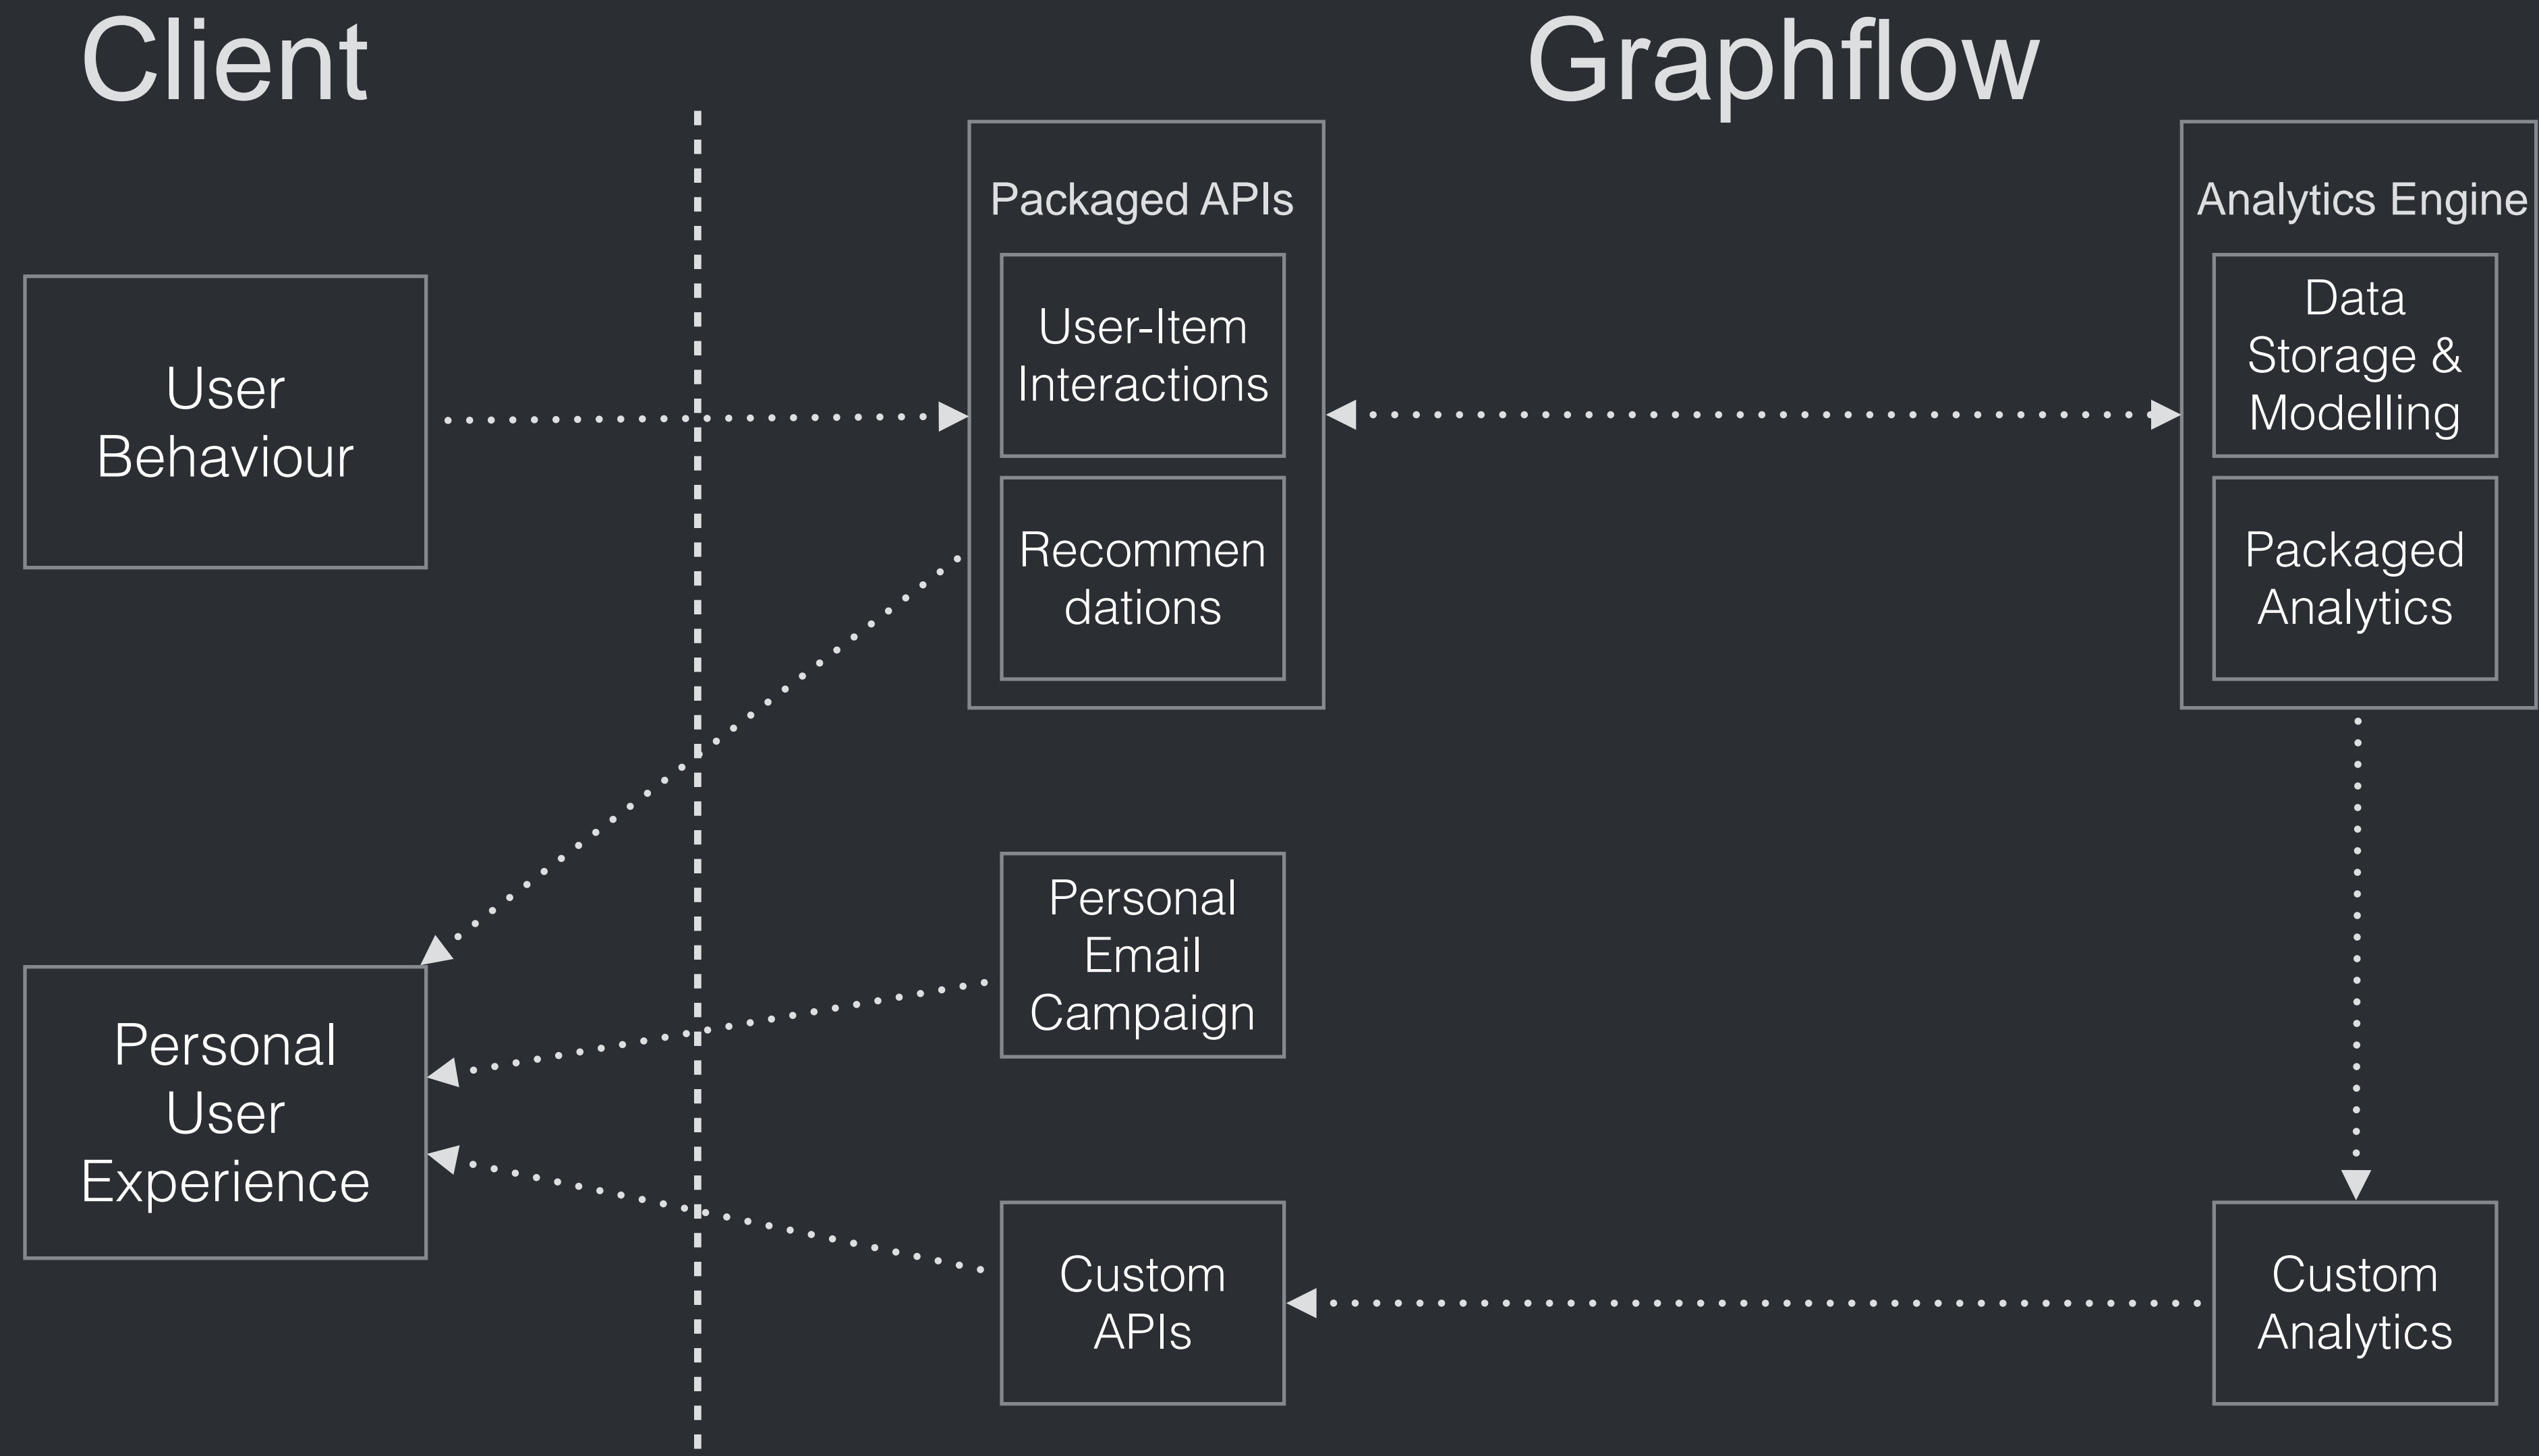
Omni-channel
Recommendations



Customer
Insight



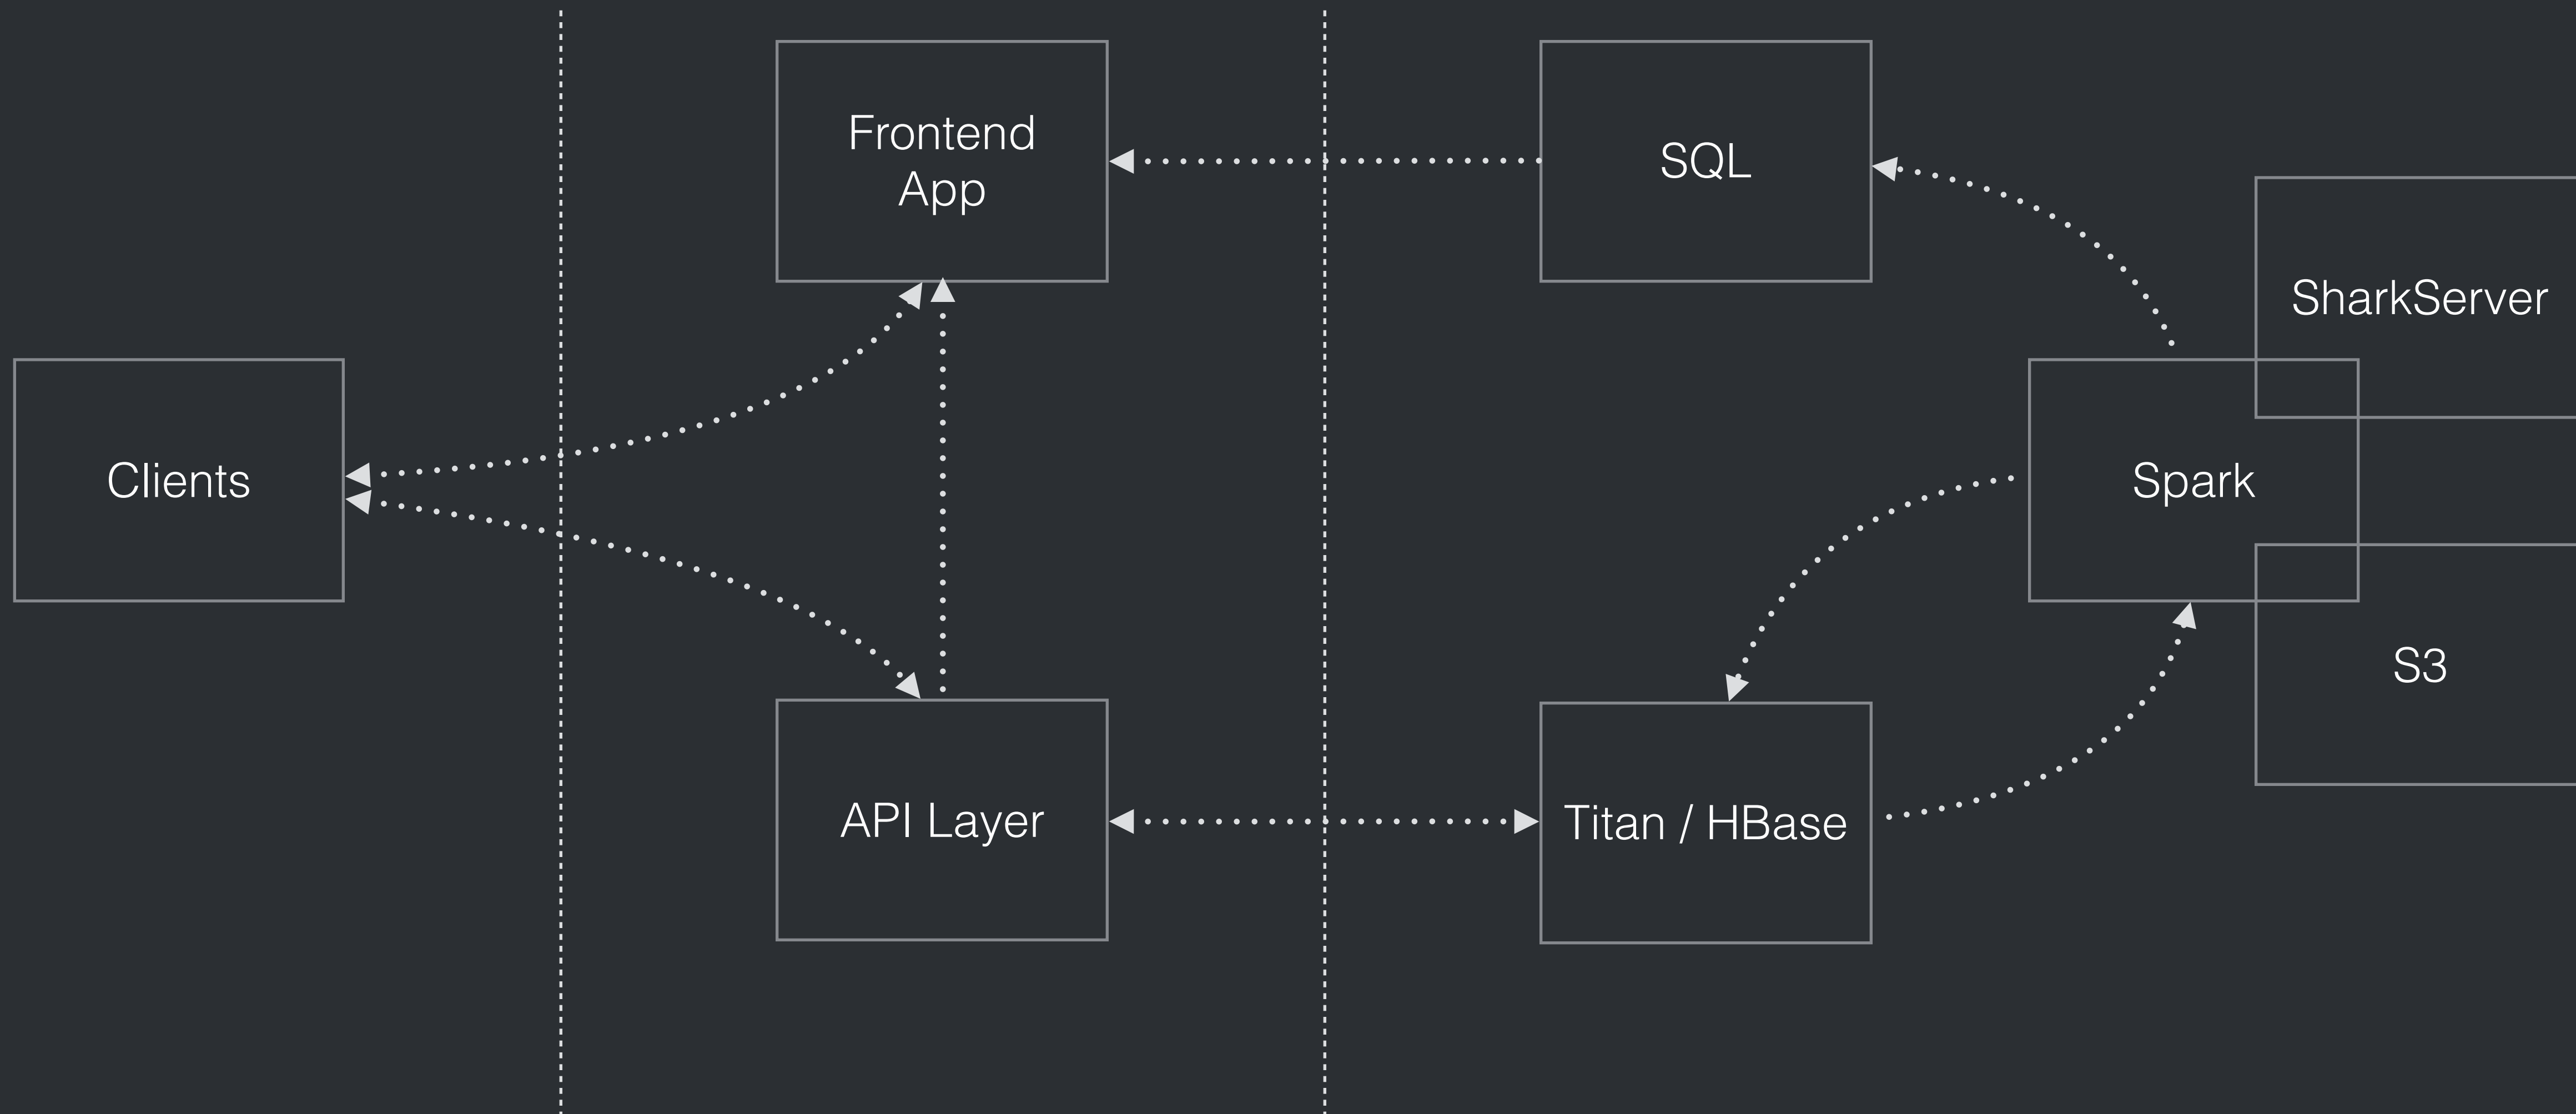
The Graphflow Platform



Why Spark?

- In-memory caching + broadcast variables == machine learning
- Scala + Python APIs + SQL (Shark)
- Spark Streaming
- (more recently)
 - *MLlib*
 - *SparkSQL*

Architecture, Take One



Titan - Spark Integration

- Faunus: Gremlin DSL on Hadoop (in Groovy)

```
g = FaunusFactory.open(conf)
g.E.has("label", "purchase").transform(...)
```

- Spark

```
val graph = sc.newAPIHadoopRDD(
    faunusGraph.getConf,
    fClass = classOf[TitanHBaseInputFormat],
    kClass = classOf[NullWritable],
    vClass = classOf[FaunusVertex])
graph.flatMap { vertex =>
    val edges = vertex.getEdges(Direction.OUT).filter(e => e.getLabel == "purchase")
    edges.map { edge => (...) }
}
```

Spark Challenges, mid-2013

- Spark 0.7.x
- No MLlib ...
 - *... so we wrote our own*
 - *Port of Mahout's Alternating Least Squares using Spark + Breeze*
- No JobServer ...
 - *... so we wrote our own*
 - *Scalatra REST API + Akka Actors*
 - *Simpler version of Ooyala's JobServer*

Spark ALS

- Spark 0.8.0-SNAPSHOT
- MovieLens 1m dataset
- Implicit prefs, 10 iterations, 10 factors

Model	Mahout 0.8 ALS	Spark Custom ALS
Runtime	6m39s	58s
Lines of Code	$1374 + 137 + 121 + 116 =$ ~1750	~425

Spark ALS

```
public class ImplicitFeedbackAlternatingLeastSquaresSolver {

    private final int numFeatures;
    private final double alpha;
    private final double lambda;

    private final OpenIntObjectHashMap<Vector> Y;
    private final Matrix YtransposeY;

    public ImplicitFeedbackAlternatingLeastSquaresSolver(int numFeatures, double lambda, double alpha,
        OpenIntObjectHashMap<Vector> Y) {
        this.numFeatures = numFeatures;
        this.lambda = lambda;
        this.alpha = alpha;
        this.Y = Y;
        YtransposeY = getYtransposeY(Y);
    }

    public Vector solve(Vector ratings) {
        return solve(YtransposeY.plus(getYtransposeCuMinusIYPlusLambdaI(ratings)), getYtransposeCuPu(ratings));
    }

    private static Vector solve(Matrix A, Matrix y) {
        return new QRDecomposition(A).solve(y).viewColumn(0);
    }

    double confidence(double rating) {
        return 1 + alpha * ratings;
    }

    /* Y' Y */
    private Matrix getYtransposeY(OpenIntObjectHashMap<Vector> Y) {

        IntArrayList indexes = Y.keys();
        indexes.quickSort();
        int numIndexes = indexes.size();

        double[][] YtY = new double[numFeatures][numFeatures];

        // Compute Y'Y by dot products between the 'columns' of Y
        for (int i = 0; i < numFeatures; i++) {
            for (int j = i; j < numFeatures; j++) {
                double dot = 0;
                for (int k = 0; k < numIndexes; k++) {
                    Vector row = Y.get(indexes.getQuick(k));
                    dot += row.getQuick(i) * row.getQuick(j);
                }
                YtY[i][j] = dot;
                if (i != j) {
                    YtY[j][i] = dot;
                }
            }
        }
        return new DenseMatrix(YtY, true);
    }

    /** Y' (Cu - I) Y + λ I */
    private Matrix getYtransposeCuMinusIYPlusLambdaI(Vector userRatings) {
        Preconditions.checkArgument(userRatings.isSequentialAccess(), "need sequential access to ratings!");

        /* (Cu -I) Y */
        OpenIntObjectHashMap<Vector> CuMinusIY = new OpenIntObjectHashMap<Vector>(userRatings.getNumNondefaultElements());
        for (Element e : userRatings.nonZeroes()) {
            CuMinusIY.put(e.index(), Y.get(e.index()).times(confidence(e.get()) - 1));
        }

        Matrix YtransposeCuMinusIY = new DenseMatrix(numFeatures, numFeatures);

        /* Y' (Cu -I) Y by outer products */
        for (Element e : userRatings.nonZeroes()) {
            for (Vector.Element feature : Y.get(e.index()).all()) {
                Vector partial = CuMinusIY.get(e.index()).times(feature.get());
                YtransposeCuMinusIY.viewRow(feature.index()).assign(partial, Functions.PLUS);
            }
        }

        /* Y' (Cu - I) Y + λ I add lambda on the diagonal */
        for (int feature = 0; feature < numFeatures; feature++) {
            YtransposeCuMinusIY.setQuick(feature, feature, YtransposeCuMinusIY.getQuick(feature, feature) + lambda);
        }

        return YtransposeCuMinusIY;
    }

    /** Y' Cu p(u) */
    private Matrix getYtransposeCuPu(Vector userRatings) {
        Preconditions.checkArgument(userRatings.isSequentialAccess(), "need sequential access to ratings!");

        Vector YtransposeCuPu = new DenseVector(numFeatures);

        for (Element e : userRatings.nonZeroes()) {
            YtransposeCuPu.assign(Y.get(e.index()).times(confidence(e.get())) , Functions.PLUS);
        }

        return columnVectorAsMatrix(YtransposeCuPu);
    }

    private Matrix columnVectorAsMatrix(Vector v) {
        double[][] matrix = new double[numFeatures][1];
        for (Vector.Element e : v.all()) {
            matrix[e.index()][0] = e.get();
        }
        return new DenseMatrix(matrix, true);
    }
}
```

VS

```
def updateFactorsImplicit(
  UorI: mutable.Map[Int, DenseVector[Double]],
  ratings: Vector[Double],
  YtY: DenseMatrix[Double]) = {

  // set up required intermediate data structures
  val nui = ratings.activeSize
  val UorIMat = DenseMatrix.zeros[Double](nui, numF)
  val CuMinusIY = DenseMatrix.zeros[Double](nui, numF)
  val Cup = DenseVector.zeros[Double](nui)
  var j = 0
```

```
  ratings.activeIterator.foreach{ case(i, v) => {
    CuMinusIY(j, ::) := UorI(i) :* alpha :* v
    Cup(j) = alpha * v + 1
    UorIMat(j, ::) := UorI(i)
    j += 1
  }}
```

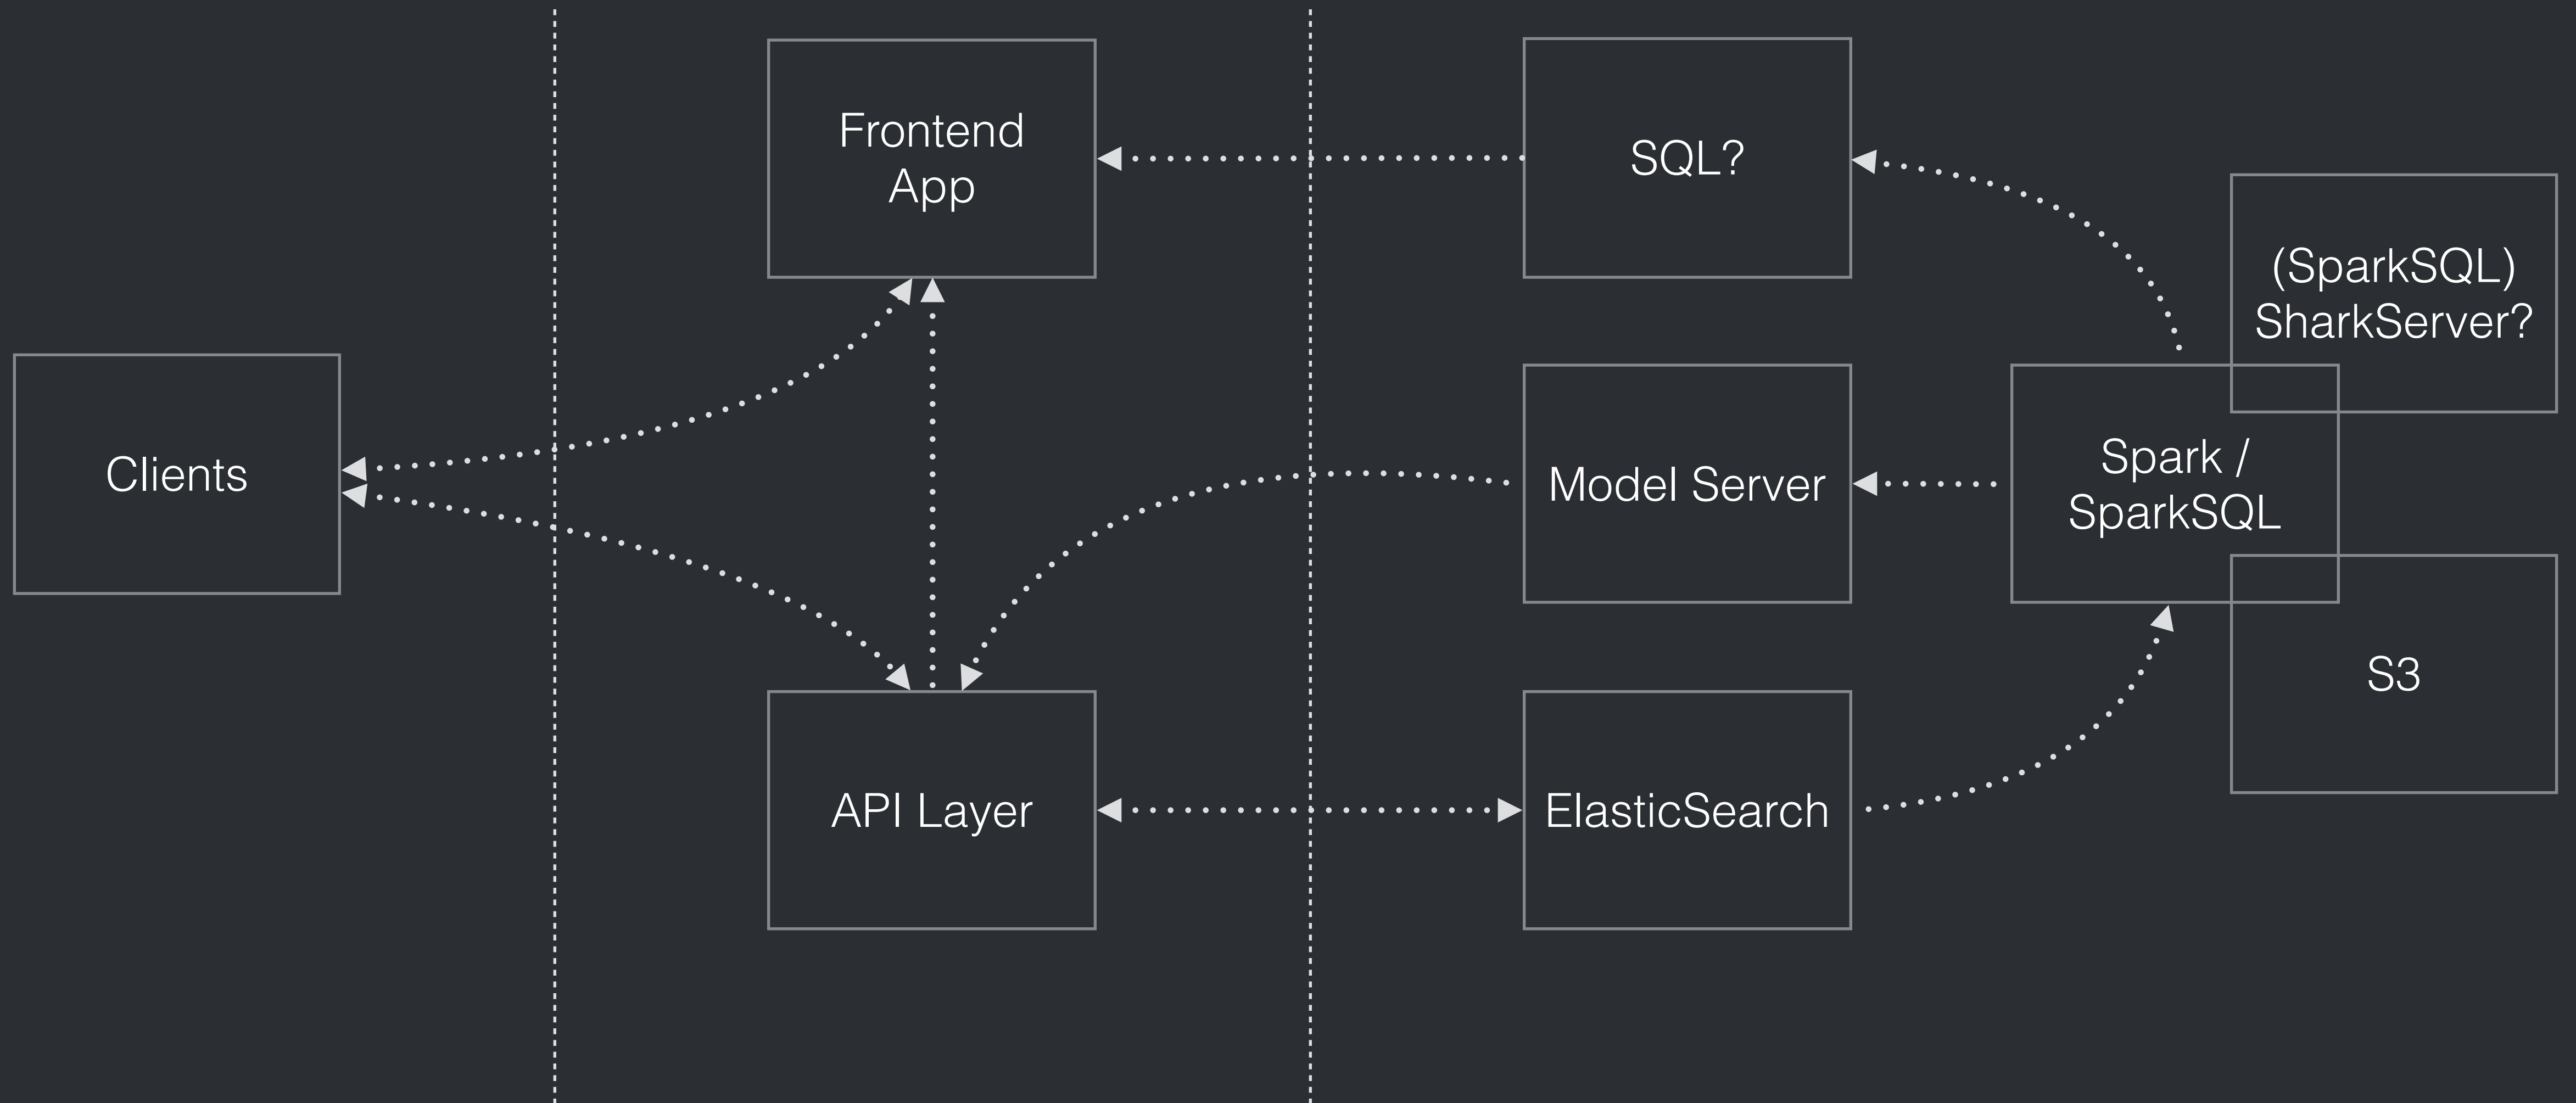
```
  val YtCuY =
    YtY + UorIMat.t * CuMinusIY + (DenseMatrix.eye[Double](numF) :* lambda)
  val YtCup = UorIMat.t * Cup
  YtCuY \ YtCup
}
```

Matrix / vector multiplication
Element-wise operations

Architecture Issues

- Ended up not using Titan's powerful features
- Search, filtering, aggregation more important
- Titan/HBase + ElasticSearch
 - *Complexity*
 - *Data duplication*
- Always aimed for real-time model serving + incremental updates

Architecture, Take Two



ElasticSearch - Spark Integration

- Faunus / Spark

```
val graph = sc.newAPIHadoopRDD(faunusGraph.getConf,  
                                classOf[TitanHBaseInputFormat],  
                                classOf[NullWritable],  
                                classOf[FaunusVertex])  
  
val events = graph.flatMap { vertex =>  
    vertex.getEdges(Direction.OUT).map { edge => (...) }  
}
```

- ElasticSearch / Spark

```
conf.set("es.resource", s"$index/$doc")  
conf.set("es.query", s"?q=$q")  
val esRdd = sc.newAPIHadoopRDD(conf,  
                                classOf[EsInputFormat[NullWritable, LinkedMapWritable]],  
                                classOf[NullWritable],  
                                classOf[LinkedMapWritable])  
  
val events = esRdd.map { case (_, m) => (...) }
```

SparkSQL

- ElasticSearch / SparkSQL

```
val events = esRdd.map { case (_, m) =>
    Event(m("timestamp").toLong, m("event").toString, ...)
}

// complex join and filter in Spark
...

events.registerAsTable("events")
val aggs = hql(
    "select
        from_unixtime(cast(time/1000.0 as bigint), 'yyyy-MM-dd HH:00:00') hour,
        event,
        count(1)
    from events ..."
)

// save results to S3, ElasticSearch, MySQL, etc...
```


Spark ALS, Take Two

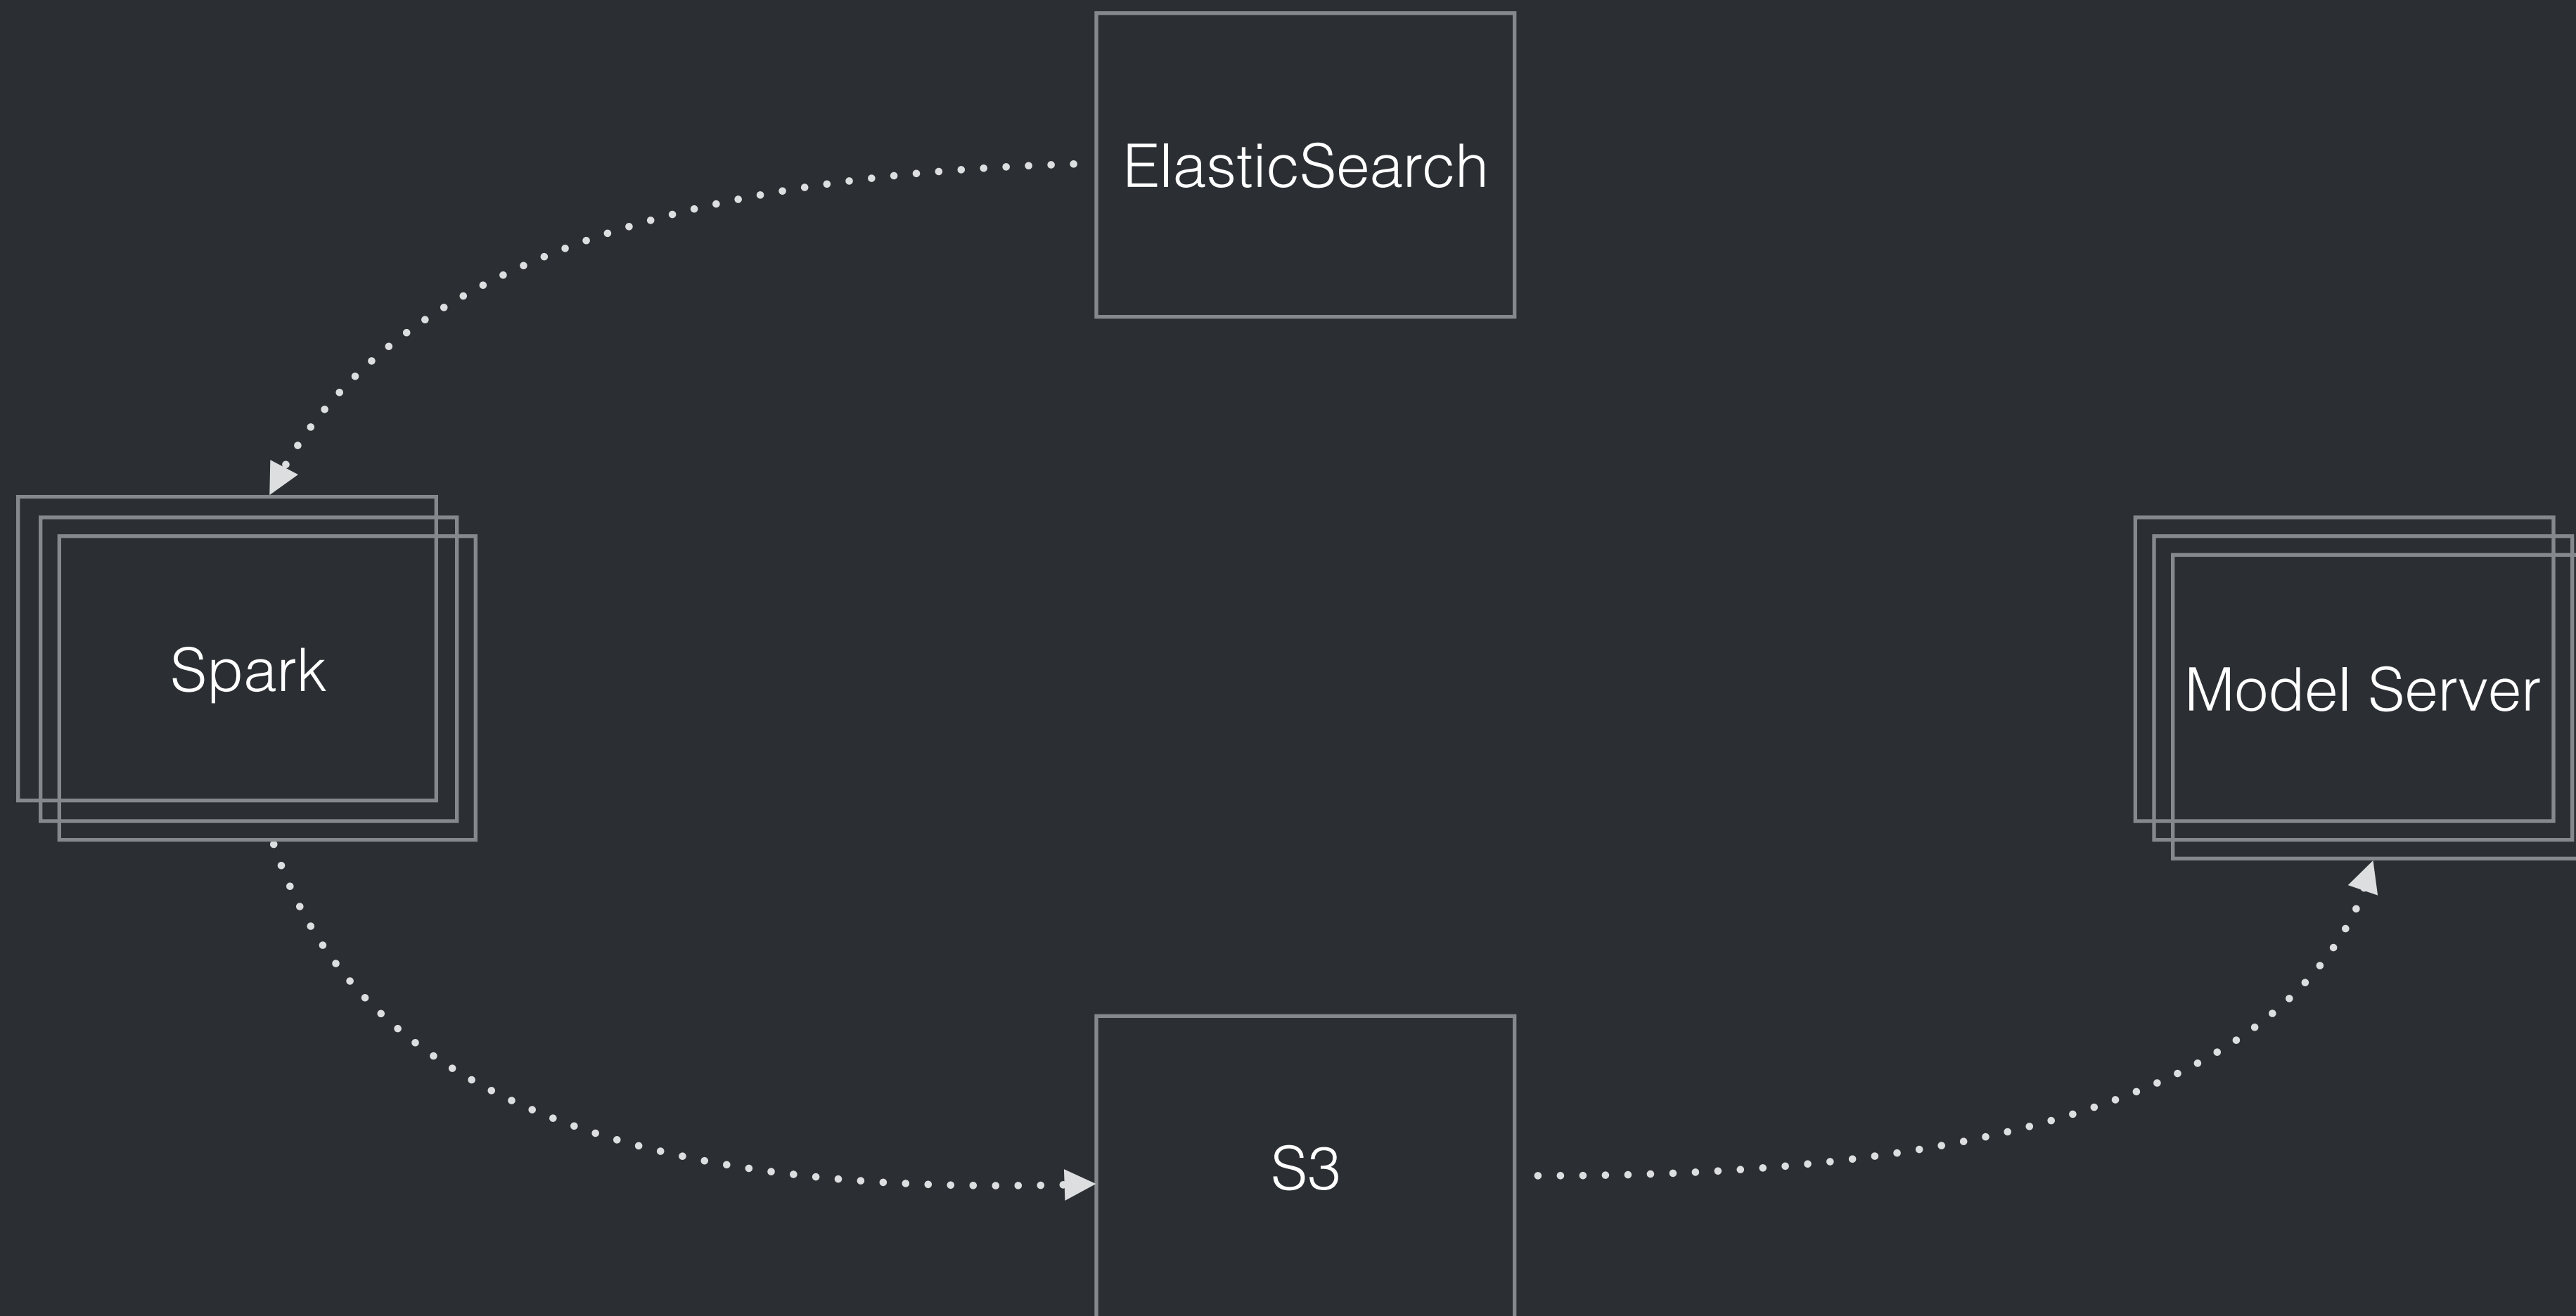
- Spark MLlib
- MovieLens 1m dataset
- Implicit prefs, 10 iterations, 10 factors

Model	Mahout 0.8 ALS	Spark Custom ALS	MLlib ALS
Runtime	6m39s	58s	24.8s
Lines of Code	$1374 + 137 + 121 + 116 =$ ~1750	~425	~880

Model Server

- Requirements:
 - *Serve 100s - 1000s concurrent models*
 - *Fast & horizontally scalable*
 - *Many small models (10s - 100s thousands users/items)*
 - *Some very large models (millions of users/items)*
 - *Fault tolerant*
- Available alternatives didn't fit: Oryx, Prediction.io

Model Server



Model Server

- Scalatra
- Akka Cluster
- Breeze
- Periodic model refresh
- Incremental updates




Spark Challenges, mid-2014

- Spark 1.0.0!
 - *MLlib*
 - *JobServer*
 - *Spark Submit*
- Main challenges now:
 - *DevOps*
 - *Deployment*
 - *Automation & scheduling*

Future

- Near future:
 - *More complex recommendation models in Spark*
 - *Different types of predictive models*
 - *More complex and compute-intensive aggregations for reporting & insights*
- Farther future:
 - *Spark Streaming for online model updates and aggregations*



Graph Flow

